

PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Update all packages with `sudo yum update -y`.
- Install Java with `sudo yum install java-11-openjdk-devel -y` or `sudo yum install java-1.8.0-amazon-corretto-devel -y`.
- Verify the Java installation using `java -version`.
- Navigate to `/usr/local/` using `cd /usr/local/`.
- Download Hadoop with `sudo wget https://downloads.apache.org/hadoop/common/hadoop-3.3.6/hadoop-3.3.6.tar.gz`.
- Extract the Hadoop archive using `sudo tar -xvzf hadoop-3.3.6.tar.gz`.
- Rename the extracted folder with `sudo mv hadoop-3.3.6 hadoop`.
- Edit the `.bashrc` file to add Hadoop and Java environment variables.
- Reload the `.bashrc` file with `source ~/.bashrc` and verify Hadoop installation using `hadoop version`.

CODING: -

- `sudo su`
- `sudo yum update -y`
- `sudo yum install java-11-openjdk-devel -y` **or** `sudo yum install java-1.8.0-amazon-corretto-devel -y`
- `java -version`
- `cd /usr/local/`
- `sudo wget https://downloads.apache.org/hadoop/common/hadoop-3.3.6/hadoop-3.3.6.tar.gz`
- `sudo tar -xvzf hadoop-3.3.6.tar.gz`
- `sudo mv hadoop-3.3.6 hadoop`
- `sudo nano ~/.bashrc`
 - `# Hadoop variables`
 - `export HADOOP_HOME=/usr/local/hadoop`
 - `export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin`
 - `# Java variables`
 - `export JAVA_HOME=$(readlink -f /usr/bin/java | sed "s:/bin/java::")`
 - `export PATH=$PATH:$JAVA_HOME/bin`
- `source ~/.bashrc`
- `hadoop version`

OUTPUT: -

```
[root@ip-172-31-46-39 local]# source ~/.bashrc
[root@ip-172-31-46-39 local]# hadoop version
Hadoop 3.3.6
Source code repository https://github.com/apache/hadoop.git -r 1be78238728da9266a4f88195058f08fd012bf9c
Compiled by ubuntu on 2023-06-18T08:22Z
Compiled on platform linux-x86_64
Compiled with protoc 3.7.1
From source with checksum 5652179ad55f76cb287d9c633bb53bbd
This command was run using /usr/local/hadoop/share/hadoop/common/hadoop-common-3.3.6.jar
[root@ip-172-31-46-39 local]# |
```

PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Update all packages with `sudo yum update -y`.
- Install Java with `sudo yum install java-11-openjdk-devel -y` or `sudo yum install java-1.8.0-amazon-corretto-devel -y`.
- Verify the Java installation using `java -version`.
- Navigate to `/usr/local/` using `cd /usr/local/`.
- Download Hadoop with `sudo wget https://downloads.apache.org/hadoop/common/hadoop-3.3.6/hadoop-3.3.6.tar.gz`.
- Extract the Hadoop archive using `sudo tar -xvzf hadoop-3.3.6.tar.gz`.
- Rename the extracted folder with `sudo mv hadoop-3.3.6 hadoop`.
- Edit the `.bashrc` file to add Hadoop and Java environment variables.
- Reload the `.bashrc` file with `source ~/.bashrc` and verify Hadoop installation using `hadoop version`.
- Edit `core-site.xml` to set the default file system and temporary directory.
- Configure `hdfs-site.xml` to set replication factor, NameNode, and DataNode directories.
- Create `mapred-site.xml` and configure the MapReduce framework to use YARN.
- Edit `yarn-site.xml` to set up ResourceManager and NodeManager services.
- Create the necessary Hadoop directories for the NameNode, DataNode, and temporary storage.
- Format the Hadoop HDFS filesystem using `hdfs namenode -format`.
- Start the HDFS services with `start-dfs.sh`.
- Start YARN services with `start-yarn.sh`.
- Check running processes with `jps` to ensure services are active.
- Access Hadoop web interfaces: NameNode at `http://localhost:9870/` and ResourceManager at `http://localhost:8088/`.

CODING: -

- `sudo su`
- `sudo yum update -y`
- `sudo yum install java-11-openjdk-devel -y` **or** `sudo yum install java-1.8.0-amazon-corretto-devel -y`
- `java -version`
- `cd /usr/local/`
- `sudo wget https://downloads.apache.org/hadoop/common/hadoop-3.3.6/hadoop-3.3.6.tar.gz`
- `sudo tar -xvzf hadoop-3.3.6.tar.gz`
- `sudo mv hadoop-3.3.6 hadoop`

- `sudo nano ~/.bashrc`
 - `# Hadoop variables`
 - `export HADOOP_HOME=/usr/local/hadoop`
 - `export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin`
 - `# Java variables`
 - `export JAVA_HOME=$(readlink -f /usr/bin/java | sed "s:/bin/java::")`
 - `export PATH=$PATH:$JAVA_HOME/bin`
- `source ~/.bashrc`
- `hadoop version`
- `Configure core-site.xml`
 - `sudo nano $HADOOP_HOME/etc/hadoop/core-site.xml`
 - `<configuration>`
 - `<property>`
 - `<name>fs.defaultFS</name>`
 - `<value>hdfs://localhost:9000</value>`
 - `</property>`
 - `<property>`
 - `<name>hadoop.tmp.dir</name>`
 - `<value>/usr/local/hadoop/tmp</value>`
 - `<description>Temporary directory for Hadoop</description>`
 - `</property>`
 - `</configuration>`
- `Configure hdfs-site.xml`
 - `sudo nano $HADOOP_HOME/etc/hadoop/hdfs-site.xml`
 - `<configuration>`
 - `<property>`
 - `<name>dfs.replication</name>`
 - `<value>1</value> <!-- Since this is a single-node setup -->`
 - `</property>`

- `<property>`
- `<name>dfs.namenode.name.dir</name>`
- `<value>file:///usr/local/hadoop/hdfs/namenode</value>`
- `</property>`
- `<property>`
- `<name>dfs.datanode.data.dir</name>`
- `<value>file:///usr/local/hadoop/hdfs/datanode</value>`
- `</property>`
- `</configuration>`
- Create the mapred-site.xml[If the mapred-site.xml.template is not present]
 - `sudo nano /usr/local/hadoop/etc/hadoop/mapred-site.xml`
 - `<configuration>`
 - `<property>`
 - `<name>mapreduce.framework.name</name>`
 - `<value>yarn</value>`
 - `</property>`
 - `</configuration>`
- Configure yarn-site.xml
 - `sudo nano $HADOOP_HOME/etc/hadoop/yarn-site.xml`
 - `<configuration>`
 - `<property>`
 - `<name>yarn.nodemanager.aux-services</name>`
 - `<value>mapreduce_shuffle</value>`
 - `</property>`
 - `<property>`
 - `<name>yarn.resourcemanager.resource-tracker.address</name>`
 - `<value>localhost:8025</value>`
 - `</property>`
 - `<property>`

- `<name>yarn.resourcemanager.scheduler.address</name>`
- `<value>localhost:8030</value>`
- `</property>`
- `<property>`
- `<name>yarn.resourcemanager.address</name>`
- `<value>localhost:8050</value>`
- `</property>`
- `</configuration>`
- Set Up Hadoop Directories
 - `sudo mkdir -p /usr/local/hadoop/hdfs/namenode`
 - `sudo mkdir -p /usr/local/hadoop/hdfs/datanode`
 - `sudo mkdir -p /usr/local/hadoop/tmp`
- Format the HDFS Filesystem
 - `hdfs namenode -format`
- Start Hadoop Services(As Non - Root User)
 - `start-dfs.sh`
 - `start-yarn.sh`
 - `Jps`
- Access the Hadoop Web Interfaces
 - **NameNode:** `http://localhost:9870/` (shows the HDFS overview)
 - **ResourceManager:** `http://localhost:8088/` (shows the YARN overview)
 - **LocalHost - 127.0.0.1 or public DNS :- 3.117.182.16**

OUTPUT: -

```
aws Services Search [Alt+S]
GNU nano 5.8 /root/.bashrc
# .bashrc

# Source global definitions
if [ -f /etc/bashrc ]; then
    . /etc/bashrc
fi

# User specific environment
if [ [ "$PATH" != "$HOME/.local/bin:$HOME/bin:" ] ]
then
    PATH="$HOME/.local/bin:$HOME/bin:$PATH"
fi
export PATH

# Uncomment the following line if you don't like systemctl's auto-paging feature:
# export SYSTEMD_PAGER=

# User specific aliases and functions

alias rm='rm -i'
alias cp='cp -i'
alias mv='mv -i'

[ Read 31 lines ]
Help Write Out Where Is Cut Execute Location Undo Set Mark To Bracket Previous
Exit Read File Replace Paste Justify Go To Line Redo Copy Where Was Next
```

```
aws Services Search [Alt+S]

We trust you have received the usual lecture from the local System
Administrator. It usually boils down to these three things:

#1) Respect the privacy of others.
#2) Think before you type.
#3) With great power comes great responsibility.

For security reasons, the password you type will not be visible.

[sudo] password for hadoop:
hadoop is not in the sudoers file.
[hadoop@ip-172-31-40-166 ~]$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [ip-172-31-40-166.ec2.internal]
[hadoop@ip-172-31-40-166 ~]$ start-yarn.sh
Starting resourcemanager
resourcemanager is running as process 33674. Stop it first and ensure /tmp/hadoop-hadoop-resourcemanager.pid file is empty before retry.
Starting nodemanagers
[hadoop@ip-172-31-40-166 ~]$ jps
35073 Jps
34913 NodeManager
34360 DataNode
34617 SecondaryNameNode
34250 NameNode
33674 ResourceManager
[hadoop@ip-172-31-40-166 ~]$

i-015e997a4b9339fce (BDA-Lab)
PublicIPs: 3.88.65.175 PrivateIPs: 172.31.40.166
```

```
# Therefore, the vast majority (BUT NOT ALL!) of these defaults
# are configured for substitution and not append. If append
# is preferable, modify this file accordingly.

###
# Generic settings for HADOOP
###

# Technically, the only required environment variable is JAVA_HOME.
# All others are optional. However, the defaults are probably not
# preferred. Many sites configure these options outside of Hadoop,
# such as in /etc/profile.d

# The java implementation to use. By default, this environment
# variable is REQUIRED on ALL platforms except OS X!
# export JAVA_HOME=$(readlink -f /usr/bin/java | sed "s:/bin/java:/" )

# Location of Hadoop. By default, Hadoop will attempt to determine
# this location based upon its execution path.
```

PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Adding Files to HDFS
 - To upload a local file to HDFS, use the `hdfs dfs -put`
- Listing Files and Directories
 - To list files and directories in a specific HDFS directory, use the `hdfs dfs -ls` command
- Creating Directories in HDFS
 - To create a new directory in HDFS, use the `hdfs dfs -mkdir` command
- Retrieving Files from HDFS
 - To download a file from HDFS to your local filesystem, use the `hdfs dfs -get` command
- Deleting Files from HDFS
 - To delete a file from HDFS, use the `hdfs dfs -rm` command
- Deleting Directories from HDFS
 - To delete a directory and its contents from HDFS, use the `hdfs dfs -rm -r` command
- Viewing File Contents in HDFS
 - To view the contents of a file in HDFS, use the `hdfs dfs -cat` command
- Checking Disk Usage of HDFS Directory
 - To check disk usage of files and directories in HDFS, use the `hdfs dfs -du -h` command
- Copying Files within HDFS
 - To copy a file or directory within HDFS, use the `hdfs dfs -cp` command
- Moving Files within HDFS
 - To move a file or directory within HDFS, use the `hdfs dfs -mv` command

CODING: -

- `sudo su`
- `hdfs dfs -put /local/path/to/file /hdfs/path/`
- `hdfs dfs -ls /hdfs/path/`
- `hdfs dfs -mkdir /hdfs/path/`
- `hdfs dfs -get /hdfs/path/to/file /local/path/`
- `hdfs dfs -rm /hdfs/path/to/file`
- `hdfs dfs -rm -r /hdfs/path/to/directory`
- `hdfs dfs -cat /hdfs/path/to/file`
- `hdfs dfs -du -h /hdfs/path/`

- `hdfs dfs -cp /hdfs/source/path /hdfs/destination/path`
- `hdfs dfs -mv /hdfs/source/path /hdfs/destination/path`

OUTPUT: -

```
aws Services Search [Alt+S]
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.1.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.2.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.3.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.8.2.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.3.3.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Null.java
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.8.3.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.3.5.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.8.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.3.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/hadoop-hdfs-0.22.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.9.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.1.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/hadoop-hdfs-0.20.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.0-alpha4.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.9.2.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.0-alpha2.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.2.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.10.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.1.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.4.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/hadoop-hdfs-0.21.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.1.3.xml
hadoop-3.3.6/share/hadoop/hdfs/hadoop-hdfs-client-3.3.6-tests.jar
hadoop-3.3.6/share/hadoop/hdfs/hadoop-hdfs-httplibs-3.3.6.jar
[root@ip-172-31-40-166 local]# sudo mv hadoop-3.3.6 hadoop
[root@ip-172-31-40-166 local]# sudo nano ~/.bashrc
[root@ip-172-31-40-166 local]# sudo nano ~/.bashrc
[root@ip-172-31-40-166 local]# source ~/.bashrc
[root@ip-172-31-40-166 local]# hadoop version
Hadoop 3.3.6
Source code repository https://github.com/apache/hadoop.git -r 1be78238728da9266a4f8b195058f08fd012bf9c
Compiled by ubuntu on 2023-06-18T08:22Z
Compiled on platform linux-x86_64
Compiled with protoc 3.7.1
From source with checksum 5652179ad5576cb287d9c633bb53bbd
This command was run using /usr/local/hadoop/share/hadoop/common/hadoop-common-3.3.6.jar
[root@ip-172-31-40-166 local]#
```

```
Bytes Read=173
File Output Format Counters
Bytes Written=200
[root@ip-172-31-40-166 local]# cat /usr/local/hadoop/output/*
For 1
Hadoop, 1
about 1
and 1
at: 2
http://hadoop.apache.org/ 1
https://cwiki.apache.org/confluence/display/HADOOP/ 1
information 1
latest 1
our 2
please 1
the 1
visit 1
website 1
wiki, 1
[root@ip-172-31-40-166 local]#
```

PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Input Format:
 - We'll have two input files representing matrices A and B.
 - Matrix A (m x n): Split by rows.
 - Matrix B (n x p): Split by columns.
- Map Step:
 - For each element in matrix A (i, k, A[i][k]), emit the product to the intermediate key (i, j) where j is the column index in matrix B.
 - For each element in matrix B (k, j, B[k][j]), emit the product to the intermediate key (i, j) for all rows i of matrix A.
- Reduce Step:
 - For each intermediate key (i, j), sum the products of the corresponding values to calculate the result matrix C (i, j).
- Output Format:
 - Output matrix C, where each line is of the form i, j, C[i][j].

CODING: -

- `sudo su`
- `MatrixMapper.java`
 - `import java.io.IOException;`
 - `import org.apache.hadoop.io.IntWritable;`
 - `import org.apache.hadoop.io.Text;`
 - `import org.apache.hadoop.mapreduce.Mapper;`
 -
 - `public class MatrixMapper extends Mapper<Object, Text, Text, Text> {`
 -
 - `@Override`
 - `public void map(Object key, Text value, Context context) throws IOException, InterruptedException {`
 - `String[] line = value.toString().split(",");`
 - `String matrixName = line[0]; // A or B`
 - `int i = Integer.parseInt(line[1]);`
 - `int j = Integer.parseInt(line[2]);`
 - `int valueOfElement = Integer.parseInt(line[3]);`

-
- if (matrixName.equals("A")) {
- // Emit for all columns of B
- for (int k = 0; k < context.getConfiguration().getInt("p", 0); k++) {
- context.write(new Text(i + "," + k), new Text("A," + j + "," + valueOfElement));
- }
- } else {
- // Emit for all rows of A
- for (int k = 0; k < context.getConfiguration().getInt("m", 0); k++) {
- context.write(new Text(k + "," + j), new Text("B," + i + "," + valueOfElement));
- }
- }
- }
- }
- }

- MatrixReducer.java

- import java.io.IOException;
- import java.util.HashMap;
- import java.util.Map;
- import org.apache.hadoop.io.IntWritable;
- import org.apache.hadoop.io.Text;
- import org.apache.hadoop.mapreduce.Reducer;
-
- public class MatrixReducer extends Reducer<Text, Text, Text, IntWritable> {
-
- @Override
- public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {
- Map<Integer, Integer> mapA = new HashMap<>();

- Map<Integer, Integer> mapB = new HashMap<>();
-
- for (Text val : values) {
- String[] parts = val.toString().split(",");
- if (parts[0].equals("A")) {
- mapA.put(Integer.parseInt(parts[1]), Integer.parseInt(parts[2]));
- } else {
- mapB.put(Integer.parseInt(parts[1]), Integer.parseInt(parts[2]));
- }
- }
-
- int result = 0;
- for (Integer k : mapA.keySet()) {
- if (mapB.containsKey(k)) {
- result += mapA.get(k) * mapB.get(k);
- }
- }
-
- context.write(key, new IntWritable(result));
- }
- }

- MatrixMultiplication.java

- import org.apache.hadoop.conf.Configuration;
- import org.apache.hadoop.fs.Path;
- import org.apache.hadoop.io.IntWritable;
- import org.apache.hadoop.io.Text;
- import org.apache.hadoop.mapreduce.Job;
- import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
- import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

```

○ public class MatrixMultiplication {
○
○     public static void main(String[] args) throws Exception {
○         Configuration conf = new Configuration();
○         // Dimensions of the matrices A (m x n) and B (n x p)
○         conf.setInt("m", 3); // Rows of A
○         conf.setInt("n", 2); // Columns of A and Rows of B
○         conf.setInt("p", 3); // Columns of B
○
○
○         Job job = Job.getInstance(conf, "Matrix Multiplication");
○         job.setJarByClass(MatrixMultiplication.class);
○         job.setMapperClass(MatrixMapper.class);
○         job.setReducerClass(MatrixReducer.class);
○
○
○         job.setOutputKeyClass(Text.class);
○         job.setOutputValueClass(IntWritable.class);
○
○
○         FileInputFormat.addInputPath(job, new Path(args[0]));
○         FileOutputFormat.setOutputPath(job, new Path(args[1]));
○         System.exit(job.waitForCompletion(true) ? 0 : 1);
○     }
○ }

```

- Steps to Run the Code

```

○ hdfs dfs -put matrixA.txt /input/
○ hdfs dfs -put matrixB.txt /input/
○ hadoop com.sun.tools.javac.Main MatrixMultiplication.java
○ jar cf matrixmultiplication.jar MatrixMultiplication*.class
○ hadoop jar matrixmultiplication.jar MatrixMultiplication /input/ /output/
○ hdfs dfs -cat /output/part-r-00000

```


PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Create a text file on your local filesystem that contains some sample data (e.g., `input.txt`).
- Put the file into HDFS (Hadoop Distributed File System)
- The mapper reads the input line by line, splits each line into words, and outputs each word as a key with a value of 1.
- The reducer sums the counts for each word emitted by the mapper and outputs the word along with its total count
- Compile the Java Code
- Run the MapReduce Job
- Once the job is complete, you can check the output by viewing the result file in HDFS

CODING: -

- `sudo su`
- `WordCountMapper.java`
 - `import java.io.IOException;`
 - `import org.apache.hadoop.io.IntWritable;`
 - `import org.apache.hadoop.io.LongWritable;`
 - `import org.apache.hadoop.io.Text;`
 - `import org.apache.hadoop.mapreduce.Mapper;`
 -
 - `public class WordCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {`
 -
 - `private final static IntWritable one = new IntWritable(1);`
 - `private Text word = new Text();`
 -
 - `@Override`
 - `protected void map(LongWritable key, Text value, Context context) throws`
`IOException, InterruptedException {`
 - `String[] words = value.toString().split("\\s+");`
 - `for (String str : words) {`

- word.set(str.replaceAll("[^a-zA-Z]", "").toLowerCase()); // Normalize word
- if (!word.toString().isEmpty()) {
- context.write(word, one);
- }
- }
- }
- }

- WordCountReducer.java

- import java.io.IOException;
- import org.apache.hadoop.io.IntWritable;
- import org.apache.hadoop.io.Text;
- import org.apache.hadoop.mapreduce.Reducer;
-
- public class WordCountReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
-
- @Override
- protected void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {
- int sum = 0;
- for (IntWritable val : values) {
- sum += val.get();
- }
- context.write(key, new IntWritable(sum));
- }
- }

- WordCount.java

- import org.apache.hadoop.conf.Configuration;
- import org.apache.hadoop.fs.Path;
- import org.apache.hadoop.io.IntWritable;
- import org.apache.hadoop.io.Text;
- import org.apache.hadoop.mapreduce.Job;
- import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
- import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
-
- public class WordCount {
-
- public static void main(String[] args) throws Exception {
- Configuration conf = new Configuration();
- Job job = Job.getInstance(conf, "Word Count");
-
- job.setJarByClass(WordCount.class);
- job.setMapperClass(WordCountMapper.class);
- job.setReducerClass(WordCountReducer.class);
-
- job.setOutputKeyClass(Text.class);
- job.setOutputValueClass(IntWritable.class);
-
- FileInputFormat.addInputPath(job, new Path(args[0]));
- FileOutputFormat.setOutputPath(job, new Path(args[1]));
-
- System.exit(job.waitForCompletion(true) ? 0 : 1);
- }
- }

- Steps to Run the Code

- `hadoop com.sun.tools.javac.Main WordCount.java`
- `jar cf wordcount.jar WordCount*.class`
- `hadoop jar wordcount.jar WordCount /input /output`
- `hdfs dfs -cat /output/part-r-00000`

OUTPUT: -

```
File Output Format Counters
  Bytes Written=200
[root@ip-172-31-40-166 local]# cat /usr/local/hadoop/output/*
For      1
Hadoop,  1
about    1
and       1
at:       2
http://hadoop.apache.org/      1
https://cwiki.apache.org/confluence/display/HADOOP/      1
information      1
latest  1
our      2
please  1
the      1
visit   1
website 1
wiki,    1
[root@ip-172-31-40-166 local]#
```

PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Update all system packages to ensure you have the latest versions.
- Install `wget` to download files from the internet.
- Use `wget` to download the Apache Hive binary tarball from the official website
- Extract the Hive binary tarball using the `tar` command.
- Move the extracted Hive directory to `/usr/local/hive` for proper management.
- Open the `.bashrc` file to set the necessary environment variables for Hive.
- After editing the `.bashrc` file, apply the changes by sourcing it.
- You can verify that Hive has been successfully installed by running the `hive` command.

CODING: -

- `sudo su`
- **Update the system**
 - `sudo yum update -y`
- **Install wget**
 - `sudo yum install wget -y`
- **Download Apache Hive**
 - `wget https://dlcdn.apache.org/hive/hive-3.1.3/apache-hive-3.1.3-bin.tar.gz`
- **Extract the Hive tarball**
 - `tar -xzf apache-hive-3.1.3-bin.tar.gz`
- **Move to root directory**
 - `sudo mv apache-hive-3.1.3-bin /usr/local/hive`
- **Configure Hive Environment Variables**
 - `nano ~/.bashrc`
 - Add the following in that file as #Hive Variables
 - `export HIVE_HOME=/usr/local/hive`
 - `export PATH=$PATH:$HIVE_HOME/bin`
 - `export HADOOP_HOME=/usr/local/hadoop # Adjust to your Hadoop installation path`
 - `export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin`
 - `source ~/.bashrc`

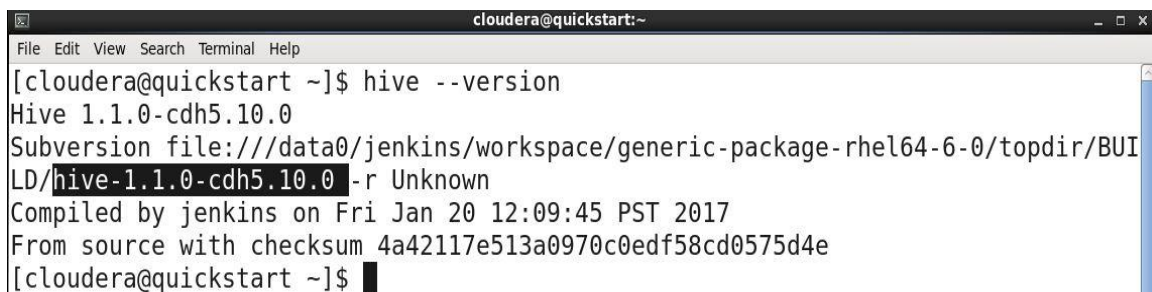
OUTPUT: -

```
# Hadoop variables
export HADOOP_HOME=/usr/local/hadoop
export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin

# Java variables
export JAVA_HOME=$(readlink -f /usr/bin/java | sed "s:/bin/java::")
export PATH=$PATH:$JAVA_HOME/bin

#Hive variables
export HIVE_HOME=/usr/local/hive
export PATH=$PATH:$HIVE_HOME/bin
export HADOOP_HOME=/usr/local/hadoop # Adjust this to match your Hadoop installation path
export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin

#MySQL variables
export PATH=$PATH:/usr/local/mysql/bin
```



A terminal window titled 'cloudera@quickstart:~' with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal shows the command '[cloudera@quickstart ~]\$ hive --version' and its output: 'Hive 1.1.0-cdh5.10.0', 'Subversion file:///data0/jenkins/workspace/generic-package-rhel64-6-0/topdir/BUILD/hive-1.1.0-cdh5.10.0 -r Unknown', 'Compiled by jenkins on Fri Jan 20 12:09:45 PST 2017', and 'From source with checksum 4a42117e513a0970c0edf58cd0575d4e'. The prompt '[cloudera@quickstart ~]\$' is shown at the bottom.

```
cloudera@quickstart:~
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ hive --version
Hive 1.1.0-cdh5.10.0
Subversion file:///data0/jenkins/workspace/generic-package-rhel64-6-0/topdir/BUILD/hive-1.1.0-cdh5.10.0 -r Unknown
Compiled by jenkins on Fri Jan 20 12:09:45 PST 2017
From source with checksum 4a42117e513a0970c0edf58cd0575d4e
[cloudera@quickstart ~]$
```

PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Update all system packages to ensure you have the latest versions.
- Install `wget` to download files from the internet.
- Use `wget` to download the HBase binary tarball from the official website.
- Extract the HBase binary tarball using the `tar` command.
- Move the extracted HBase directory to `/usr/local/hbase` for proper management.
- Open the `.bashrc` file to set the necessary environment variables for HBase.
- After editing the `.bashrc` file, apply the changes by sourcing it.
- Navigate to the HBase configuration directory.
- Open and edit the `hbase-env.sh` file to configure the environment variables for HBase.
- Open the `hbase-site.xml` file to configure HBase's storage mode.
- Finally, start the HBase services using the following command.
- You can verify if HBase has started correctly by running the HBase shell

CODING: -

- `sudo su`
- **Update the system**
 - `sudo yum update -y`
- **Install wget**
 - `sudo yum install wget -y`
- **Download HBase**
 - `wget https://downloads.apache.org/hbase/2.4.15/hbase-2.4.15-bin.tar.gz`
- **Extract the HBase tarball**
 - `tar -xzf hbase-2.4.15-bin.tar.gz`
- **Move to root directory**
 - `sudo mv hbase-2.4.15 /usr/local/hbase`
- **Configure Hbase Environment Variables**
 - `nano ~/.bashrc`
 - Add the following in that file as #HBase Variables
 - `export HBASE_HOME=/usr/local/hbase`
 - `export PATH=$PATH:$HBASE_HOME/bin`
 - `export JAVA_HOME=/usr/lib/jvm/java-1.8.0`
 - `source ~/.bashrc`
- **Configure HBase**
 - `cd /usr/local/hbase/conf`
 - `sudo nano hbase-env.sh`
 - `export JAVA_HOME=/usr/lib/jvm/java-1.8.0`
- **Configure HBase storage mode**
 - `sudo nano hbase-site.xml`
 - `<configuration>`
 - `<property>`
 - `<name>hbase.rootdir</name>`

- <value>file:///usr/local/hbase/data</value>
- </property>
- <property>
- <name>hbase.zookeeper.property.dataDir</name>
- <value>/usr/local/hbase/zookeeper</value>
- </property>
- </configuration>
- start-hbase.sh

OUTPUT: -

```
aws | Services | Search [Alt+S]
GNU nano 5.8 /root/.bashrc

# Java variables
export JAVA_HOME=$(readlink -f /usr/bin/java | sed "s:/bin/java:/:")
export PATH=$PATH:$JAVA_HOME/bin

#Hive variables
export HIVE_HOME=/usr/local/hive
export PATH=$PATH:$HIVE_HOME/bin
export HADOOP_HOME=/usr/local/hadoop # Adjust this to match your Hadoop installation path
export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin

#MySQL variables
export PATH=$PATH:/usr/local/mysql/bin

#Sqoop Variables
export SQOOP_HOME=/home/ec2-user/sqoop
export PATH=$PATH:$SQOOP_HOME/bin

#Hbase Variable
export HBASE_HOME=/usr/local/hbase
export PATH=$PATH:$HBASE_HOME/bin
export JAVA_HOME=/usr/lib/jvm/java-1.8.0

^G Help      ^O Write Out  ^W Where Is   ^K Cut        ^I Execute    ^C Location   M-U Undo
^X Exit      ^R Read File  ^N Replace    ^U Paste      ^J Justify    ^_ Go To Line  M-E Redo

ubuntu: ~/hbase-1.1.1/bin
hduser@ubuntu:~/hbase-1.1.1/bin$ hbase shell
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hduser/hbase-1.1.1/
SLF4J: Found binding in [jar:file:/home/hduser/hadoop-2.2.0
SLF4J: See http://www.slf4j.org/codes.html#multiple_binding
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLogge
2015-09-11 17:01:42,907 WARN [main] util.NativeCodeLoader:
HBase Shell; enter 'help<RETURN>' for list of supported com
Type "exit<RETURN>" to leave the HBase Shell
Version 1.1.1, rd0a115a7267f54e01c72c603ec53e91ec418292f, T
hbase(main):001:0> status
```

PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Open the HBase shell by executing the following command
- To create a table, use the `create` command
- To insert data into the table, use the `put` command.
- To retrieve data from the table, use the `get` command
- To scan data from the table, use the `scan` command
- To update a value, you can use the `put` command
- To delete a specific cell, use the `delete` command

CODING: -

- `sudo su`
- `# Start HBase shell`
- `hbase shell`
- `# Create a table`
- `create 'my_table', 'my_column_family'`
- `# Insert data (Create)`
- `put 'my_table', 'row1', 'my_column_family:name', 'Alice'`
- `put 'my_table', 'row1', 'my_column_family:age', '30'`
- `put 'my_table', 'row2', 'my_column_family:name', 'Bob'`
- `put 'my_table', 'row2', 'my_column_family:age', '25'`
- `# Read data (Retrieve)`
- `get 'my_table', 'row1'`
- `scan 'my_table'`
- `# Update data`
- `put 'my_table', 'row1', 'my_column_family:age', '31'`
- `# Delete data`
- `delete 'my_table', 'row1', 'my_column_family:age'`
- `deleteall 'my_table', 'row1' # Deletes entire row1`
- `# Drop the table (optional)`
- `drop 'my_table'`

OUTPUT: -

```
hbase(main):011:0> create 'TEST_TABLE', 'Col_1', 'Col_2'
0 row(s) in 0.3440 seconds

=> Hbase::Table - TEST_TABLE
hbase(main):012:0> list
TABLE
TEST_TABLE
1 row(s) in 0.0140 seconds

=> ["TEST_TABLE"]
```

```
ubuntu: ~/hbase-1.1.1/bin
hduser@ubuntu:~/hbase-1.1.1/bin$ hbase shell
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hduser/hbase-1.1.1/
SLF4J: Found binding in [jar:file:/home/hduser/hadoop-2.2.0
SLF4J: See http://www.slf4j.org/codes.html#multiple_binding
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLogge
2015-09-11 17:01:42,907 WARN [main] util.NativeCodeLoader:
HBase Shell; enter 'help<RETURN>' for list of supported com
Type "exit<RETURN>" to leave the HBase Shell
Version 1.1.1, rd0a115a7267f54e01c72c603ec53e91ec418292f, T
hbase(main):001:0> status
```


PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Import Data from MySQL to Hive
 - Check Sqoop Installation
 - Check MySQL Schema and Table Structure
 - Log in to MySQL to check the schema and the structure of the table you want to import
 - Describe the table to check its structure
 - Import MySQL Data into Hive using Sqoop
 - Verify the Hive Table
 - Describe the Hive table
- Export Data from Hive to MySQL
 - Check Hive Table Structure
 - Export Hive Data to MySQL with the use of sqoop commands
 - Check the Exported Data in MySQL

CODING: -

- `sudo su`
- # Part 1: Import Data from MySQL to Hive
 - # Check Sqoop Installation
 - `sqoop version`
 - # Check MySQL Schema and Table Structure
 - `mysql -u your_user -p`
 - `DESCRIBE your_database.your_table;`
 - # Import MySQL Data into Hive using Sqoop
 - `sqoop import \`
 - `--connect jdbc:mysql://your-mysql-server-ip:3306/your_database \`
 - `--username your_user \`
 - `--password your_password \`
 - `--table your_table \`
 - `--hive-import \`
 - `--hive-table your_hive_database.your_hive_table \`
 - `--m 1`

- # Verify the Hive Table
- hive
- DESCRIBE your_hive_database.your_hive_table;
- SELECT * FROM your_hive_database.your_hive_table LIMIT 10;
- # Part 2: Export Data from Hive to MySQL
 - # Check Hive Table Structure
 - DESCRIBE your_hive_database.your_hive_table;
 - # Export Hive Data to MySQL
 - sqoop export \
 - --connect jdbc:mysql://your-mysql-server-ip:3306/your_database \
 - --username your_user \
 - --password your_password \
 - --table your_table \
 - --export-dir /user/hive/warehouse/your_hive_table \
 - --m 1
 - # Check the Exported Data in MySQL
 - SELECT * FROM your_database.your_table LIMIT 10;

OUTPUT: -

```
mysql> use demo_database;  
Reading table information for completion of table and column names  
You can turn off this feature to get a quicker startup with -A
```

Database changed

```
mysql> DESCRIBE demo_table;
```

| Field | Type | Null | Key | Default | Extra |
|-------|--------------|------|-----|---------|----------------|
| id | int | NO | PRI | NULL | auto_increment |
| name | varchar(100) | NO | | NULL | |
| age | int | NO | | NULL | |

3 rows in set (0.00 sec)

Database changed

```
mysql> DESCRIBE demo_table;
```

| Field | Type | Null | Key | Default | Extra |
|-------|--------------|------|-----|---------|----------------|
| id | int | NO | PRI | NULL | auto_increment |
| name | varchar(100) | NO | | NULL | |
| age | int | NO | | NULL | |

3 rows in set (0.00 sec)

```
mysql> SELECT * FROM demo_table
```

-> ^C

```
mysql> SELECT * FROM demo_table;
```

| id | name | age |
|----|---------|-----|
| 1 | Alice | 30 |
| 2 | Bob | 25 |
| 3 | Charlie | 35 |

3 rows in set (0.00 sec)